

Zhuohan Xie, Ph.D.

Research Scientist, MBZUAI AI for Finance | Financial Agents | Trustworthy LLM Evaluation | Benchmark Design

✉ zhuohan.xie18@gmail.com 🌐 Website 🔗 LinkedIn 🐙 GitHub 🎓 Scholar

Updated September 2026

Summary: I am a Research Scientist at MBZUAI working on AI for Finance, financial agents, and trustworthy LLM evaluation. I build benchmarks, shared tasks, and verification systems for financial reasoning, evidence-grounded decision-making, and multilingual and multimodal evaluation. I also lead community evaluation initiatives and mentor early-career researchers from problem formulation through publication and release.

Selected First- and Last-Author Publications

- **[ACL 2026 Main, Oral] FinChain: A Symbolic Benchmark for Verifiable Chain-of-Thought Financial Reasoning**
Zhuohan Xie, Daniil Orel, Rushil Thareja, Dhruv Sahnan, ..., Salem Lahlou, Veselin Stoyanov, Sophia Ananiadou, Preslav Nakov
First author. Financial reasoning benchmark; symbolic verification for chain-of-thought evaluation.
- **[ACL 2026 Main, Oral] SAHM: A Benchmark for Arabic Financial and Shari'ah-Compliant Reasoning**
Rania Elbadry, ..., Sophia Ananiadou, Veselin Stoyanov, Jimin Huang, Xueqing Peng, Preslav Nakov, **Zhuohan Xie**
Last author. Arabic financial reasoning and Shari'ah-compliant evaluation.
- **[ACL 2026 Findings] FinCARDS: Card-Based Analyst Reranking for Financial Document Question Answering**
Yixi Zhou, Fan Zhang, Yu Chen, Haipeng Zhang, Preslav Nakov, **Zhuohan Xie**
Last author. Financial document QA and analyst-oriented reranking.
- **[ACL 2026 Demo] FinReporting: An Agentic Workflow for Localized Reporting of Cross-Jurisdiction Financial Disclosure**
Fan Zhang, Mingzi Song, Rania Elbadry, ..., Preslav Nakov, **Zhuohan Xie**
Last author. Agentic financial reporting system.
- **[SIGIR 2026 Oral] Multi-Sourced, Multi-Agent Evidence Retrieval for Fact-Checking**
Shuzhi Gong, Richard O. Sinnott, Jianzhong Qi, Cécile Paris, Preslav Nakov, **Zhuohan Xie**
Last author. Evidence retrieval and multi-agent fact-checking.
- **[EMNLP 2026 Findings] Right Frame, Wrong Rule: Cultural Cues Expose the Financial Knowledge Gap They Were Meant to Close**
Rania Elbadry, Ahmed Heakl, Saeed Almheiri, ..., Preslav Nakov, Yuxia Wang, Fajri Koto, **Zhuohan Xie**
Last author; accepted, forthcoming. Cultural framing and financial reasoning evaluation.
- **[AAACL 2026 Findings] SQLStructEval: Structural Evaluation of LLM Text-to-SQL Generation**
Yixi Zhou, Fan Zhang, Zhiqiao Guo, Yu Chen, Haipeng Zhang, Preslav Nakov, **Zhuohan Xie**
Last author; accepted, forthcoming. Structural evaluation of generated SQL queries.
- **[NAACL 2025 Findings] FIRE: Fact-checking with Iterative Retrieval and Verification**
Zhuohan Xie, Rui Xing, Yuxia Wang, ..., Iryna Gurevych, Preslav Nakov
First author. Fact-checking with iterative evidence retrieval and verification.
- **[EMNLP 2023 Findings] DeltaScore: Fine-Grained Story Evaluation with Perturbations**
Zhuohan Xie, Miao Li, Trevor Cohn, Jey Lau
First author. Perturbation-based evaluation of generated stories.
- **[INLG 2023] The Next Chapter: A Study of Large Language Models in Storytelling**
Zhuohan Xie, Trevor Cohn, Jey Han Lau
First author. LLM storytelling evaluation.
- **[CLEF 2026, LNCS] Overview of FinMMEval 2026: Multilingual and Multimodal Financial Evaluation**
Zhuohan Xie, Yuyang Dai, Rania Elbadry, ..., Xue Liu, Preslav Nakov
First author; forthcoming. Overview of the FinMMEval 2026 lab.
- **[arXiv 2026] FinPersona-Bench: A Benchmark for Longitudinal Psychometric Stability of Autonomous Financial Agents**
Muhammad Usman Safder, Ayesha Gull, Rania Elbadry, ..., Preslav Nakov, **Zhuohan Xie**
Last author. Long-horizon financial-agent evaluation for mandate stability and selective re-grounding.
- **[arXiv 2026] Can LLMs Be CEOs? Benchmarking Strategic Resource Reallocation with Multi-Role Agent Simulation**
Yuyang Dai, Xueqing Peng, Lingfei Qian, **Zhuohan Xie**
Last author. Constraint-rich executive decision-making with multi-role agent simulation.

- **[arXiv 2026] Seeing Is Not Deciding: Can Multimodal LLMs Act as Effective CEOs?**
Yuyang Dai, Xueqing Peng, Yuxia Wang, Preslav Nakov, **Zhuohan Xie**
Last author. Multimodal executive-decision benchmark separating perception from constrained action.
- **[arXiv 2026] Style Wins, Substance Loses: A Diagnosis of LLM-as-Judge in Idea Generation**
Fengxian Ji, Yuke Li, Jingpu Yang, . . . , Min Peng, Qianqian Xie, Xiuying Chen, **Zhuohan Xie**
Last author. Bias diagnosis and mitigation for LLM-based evaluation of scientific ideas.

Selected Additional Publications

- **[ACL 2026 Main, Oral] Is Human-Like Text Liked by Humans? Multilingual Human Detection and Preference Against AI**
Yuxia Wang, Rui Xing, . . . , **Zhuohan Xie**, . . . , Nizar Habash, Iryna Gurevych, Preslav Nakov
Co-author. Multilingual AI text detection and human preference.
- **[ACL 2026 Main] Cultural Benchmarking of LLMs in Standard and Dialectal Arabic Dialogues**
Muhammad Dehan Al Kautsar, Saeed Almheiri, . . . , **Zhuohan Xie**, Junhong Liang, Mohammad Rustom Al Nasar, Preslav Nakov, Fajri Koto
Co-author. Cultural evaluation across Standard Arabic and regional dialects.
- **[ACL 2026 Findings] Same Claim, Different Judgment: Benchmarking Scenario-Induced Bias in Multilingual Financial Misinformation Detection**
Zhiwei Liu, . . . , **Zhuohan Xie**, . . . , Alejandro Lopez-Lira, Xue Liu, Jimin Huang, Sophia Ananiadou
Co-author. Financial misinformation and scenario-induced bias.
- **[ACL 2026 Findings] RealFin: How Well Do LLMs Reason About Finance When Users Leave Things Unsaid?**
Yuyang Dai, Yan Lin, **Zhuohan Xie**, Yuxia Wang
Co-author. Implicit-assumption financial reasoning.
- **[SIGIR 2026 Resource] Conv-FinRe: A Conversational and Longitudinal Benchmark for Utility-Grounded Financial Recommendation**
Yan Wang, Yi Han, Lingfei Qian, . . . , **Zhuohan Xie**, . . . , Jimin Huang, Yankai Chen, Xue Liu, Jian-Yun Nie
Co-author. Financial recommendation benchmark.
- **[EMNLP 2026 Main] Jais 2: A Family of Arabic-Centric Open Large Language Models**
Mohamed Anwar, . . . , Yuxia Wang, **Zhuohan Xie**, Ali Mekky, . . . , Michalis Vazirgiannis, Preslav Nakov
Co-author; accepted, forthcoming. Arabic-centric open large language models.
- **[EMNLP 2026 Main] Aligning Cultural Behaviors in Arabic Language Models via Preferences**
Junhong Liang, Daniil Orel, Saeed Almheiri, . . . , Abed Alhakim Freihat, **Zhuohan Xie**, Momina Ahsan, . . . , Preslav Nakov, Fajri Koto
Co-author; accepted, forthcoming. Preference-based alignment of cultural behaviors in Arabic language models.
- **[TACL 2026] Can LLMs Automate Fact-Checking Article Writing?**
Dhruv Sahnan, . . . , **Zhuohan Xie**, Iryna Gurevych, Elizabeth Churchill, Tanmoy Chakraborty, Preslav Nakov
Co-author. Fact-checking article writing and LLM-assisted verification.
- **[ACL 2025 Main] KazMMLU: Evaluating Language Models on Kazakh, Russian, and Regional Knowledge of Kazakhstan**
Mukhammed Togmanov, Nurdaulet Mukhituly, . . . , Akhmed Sakip, **Zhuohan Xie**, Yuxia Wang, . . . , Preslav Nakov, Fajri Koto
Co-author. Evaluation of Kazakh, Russian, and regional knowledge.
- **[ACL 2025 Findings] Entity Framing and Role Portrayal in the News**
Tarek Mahmoud, **Zhuohan Xie**, Dimitar Dimitrov, . . . , Giovanni Da San Martino, Jakub Piskorski, Preslav Nakov
Co-author. Multilingual entity framing and role portrayal in news.
- **[EMNLP 2025 Main] A Head to Predict and a Head to Question: Pre-trained Uncertainty Quantification Heads for Hallucination Detection in LLM Outputs**
Artem Shelmanov, Ekaterina Fadeeva, Akim Tsvigun, Ivan Tsvigun, **Zhuohan Xie**, . . . , Preslav Nakov, Timothy Baldwin
Co-author. Pre-trained uncertainty heads for hallucination detection.
- **[COLM 2025] Sherkala-Chat: Building a State-of-the-Art LLM for Kazakh in a Moderately Resourced Setting**
Fajri Koto, Rituraj Joshi, Nurdaulet Mukhituly, Yuxia Wang, **Zhuohan Xie**, . . .
Co-author. Open Kazakh LLM with multilingual pre-training, instruction tuning, and safety alignment.
- **[FinNLP 2026, Poster] Prediction-Market Seed Capital Recovery from Noise-Dominant Flow**
Yankai Chen, Bowei He, **Zhuohan Xie**, Zhiwei Liu, Fan Zhang, Xueqing Peng, Yan Wang, Lingfei Qian, Xue Liu
Co-author; accepted Long Paper; manuscript not yet public. Prediction-market mechanism design and seed-capital recovery.

- **[SemEval 2025] SemEval 2025 Task 10: Multilingual Characterization and Extraction of Narratives from Online News**
Jakub Piskorski, Tarek Mahmoud, . . . , Nicolas Stefanovitch, **Zhuohan Xie**, Preslav Nakov, Giovanni Da San Martino
Co-author and task organizer. Shared-task overview of multilingual entity framing and narrative analysis.
- **[arXiv 2026] Herculean: An Agentic Benchmark for Financial Intelligence**
Xueqing Peng, **Zhuohan Xie**, Yupeng Cao, Haohang Li, Lingfei Qian, . . . , Sophia Ananiadou
Co-author. Large-scale benchmark for agentic financial intelligence.

Research Program

- **Vision: Reliable Evaluation Infrastructure for High-Stakes LLMs.** I build benchmarks, shared tasks, and verification systems that test whether LLMs can reason with evidence, structured constraints, multilingual inputs, and domain risk rather than simply generate plausible text.
- **Outputs: Finance, Agents, and Verification.** This program spans verifiable financial reasoning (FinChain, SAHM, FinCARDS), multilingual and multimodal evaluation (FinMMEval), agentic decision benchmarks (Herculean, FinPersona-Bench, CEO-Bench, C-SuiteBench), and evidence systems (FIRE and multi-agent retrieval).

Research Leadership and Academic Service

May 2026 – Present	Chair , Asia-Pacific Chapter, The Fin AI.
2026	Area Chair , FinNLP 2026, a workshop at EMNLP 2026.
2026–2027	Lab Organizer , CLEF FinMMEval. Organized the 2026 lab, including participant communication, evaluation workflow, and working-note coordination. The 2027 lab has been accepted and is upcoming.
2025–Present	Research Mentoring , Rania Elbadry, Yixi Zhou, Shuzhi Gong, Fan Zhang, Fengxian Ji, and Yuyang Dai; co-developed ACL/SIGIR papers and 2026 benchmark releases.
2025–2027	Organizer and Committee Member , SemEval 2025 Task 10 Organizer; SemEval 2027 Task Selection Committee.
2023–Present	Reviewer , ACL Rolling Review, NeurIPS, COLM, EMNLP, ALTA, CHOMPS, TFAI, and MisD.

Appointments and Education

May 2026 – Present	Research Scientist , MBZUAI, Abu Dhabi, UAE
Jul 2024 – May 2026	Postdoctoral Associate , MBZUAI, Abu Dhabi, UAE
Feb 2024 – Jul 2024	Research Assistant , MBZUAI, Abu Dhabi, UAE
Feb 2020 – Jul 2024	Ph.D. in Computing and Information Systems , University of Melbourne. Thesis: <i>Automatic Story Generation and Evaluation</i> .
Jul 2017 – Jun 2019	Master of Information Technology , University of Melbourne. Thesis: <i>Chinese Poetry Automatic Generation</i> .
Sep 2013 – Jun 2017	Bachelor of Software Engineering and Bachelor of Advertising , Xiamen University.

Awards and Funding

2026	Chenxing Program , CNY 300,000.
2020	Melbourne Research Scholarship , full tuition + AUD 30,600/year (2019 rate).
Honours	Dean's Honours List, School of Engineering, University of Melbourne.

Teaching and Skills

2019–2023	Academic Tutor, University of Melbourne.
Skills	Python, PyTorch, Hugging Face Transformers, OpenAI API, LangChain, scikit-learn, Bash, LaTeX; English and Mandarin Chinese.